CHAPTER 5

THE ETHICS OF CONSEQUENCES: UTILITARIANISM

Source:
Lawrence M. Hinman, *Ethics: A Pluralistic Approach to Moral Theory*. 3rd Edition. Belmont, CA: Wadsworth, 2003.

In the 1970s, Americans became increasingly aware of the carcinogenic effects of asbestos. Becoming aware of its cancer-causing consequences was no easy matter, for asbestos often does not cause cancer until twenty or more years after exposure to the asbestos dust. During World War II, a number of American factory and dock workers handled large quantities of asbestos needed for the war effort. In the 1960s, a disproportionately large percentage of them were coming down with lung cancer—and researchers began to realize that this was the result of their earlier exposure to asbestos.

By the time the harmful effects of asbestos exposure were discovered, this substance was being widely used throughout the country for insulation and brake linings and pads. Asbestos was in office buildings, school buildings, nursing homes, and private residences. Because of its harmful nature, it was difficult and costly to remove. Workers had to have special training and wear special protective clothing and breathing apparatus before they could safely work with it. One of the questions we faced as a country was precisely what to do about this problem.

To answer this question, we had to look at the consequences of either leaving the asbestos in place (perhaps with an appropriate warning) or of removing it (at high cost). When we looked at the consequences, we had to add up the potential costs and benefits of the various courses of action in various types of situations. The clearest case was in elementary schools. Children exposed to asbestos might come down with lung cancer in their thirties; in addition, many of those who worked at those schools were comparatively young. Requiring the removal of asbestos from elementary schools was an easy decision: Whatever the costs of removing it, the

potential damage was so great that few would disagree with regulations requiring its removal from schools. Should we also require its removal from all individual homes? Here the decision was more difficult, for the costs were higher proportionate to the benefits. For every elementary school in the country, there are probably hundreds of private residences insulated with asbestos. The work involved in removing asbestos from all those residences would be considerably more than the work required for the schools, and the enforcement task would be much greater as well. Moreover, the cost would presumably be borne by the individual residents or owners, who might be neither able nor willing to shoulder such a burden. Whereas a young couple with children would probably be willing to do whatever is necessary to remove the asbestos from their home, a retired couple in their eighties on a fixed income would presumably be more reluctant to spend a lot of money for this purpose. A utilitarian might well decide to leave this decision in each individual's hands. Finally, should we mandate the removal of asbestos from nursing homes and other such facilities, where the majority of people would not live an additional thirty years to be affected by asbestos? What about the welfare of those who were younger and worked in such facilities or visited them regularly? Utilitarians have to look closely at the consequences in each type of case.

Utilitarianism begins with one of the most important moral insights of modern times and couples it with a powerful metaphor which underlies our moral life. The insight is that *consequences count;* indeed, it goes one step further than this and claims that *only* consequences count. This puts it in sharp contrast to Kant's moral philosophy, which—as we shall see in the next chapter—places almost exclusive emphasis on the intentions behind an action. Utilitarianism goes to the other extreme, maintaining that **the morality of an action is to be determined solely through an assessment of its consequences.** It is for this reason that we call utilitarianism a *consequentialist* moral doctrine; morality, for the utilitarian, is solely a matter of consequences. (Utilitarianism is not the only consequentialist doctrine we have seen. Ethical egoism is also consequentialist but demands that we consider consequences only insofar as they affect our own individual well-being. Utilitarianism demands that we consider the impact of the consequences on *everyone* affected by the matter under consideration. The morally right action, the one that we ought to perform, is the one that produces the greatest overall positive consequences for everyone.

Once utilitarians have claimed that morality is solely a matter of consequences, they have to address themselves to several questions. First, they need to specify the *yardstick* or criterion in terms of which consequences are measured. Typically, utilitarians claim that we ought to do whatever produces the greatest amount of utility. But then *utility* must be defined. Pleasure, happiness, and preference satisfaction are the three most common candidates for the definition of utility. Second, utilitarians need to indicate *how the consequences can be measured.* They need, in other words, to provide an account of how the yardstick can be applied, how utility can be measured.

Third, utilitarians must address the question of how high their standards are, of *how much utility we must strive for.* How much, in other words, is enough utility? Although utilitarianism has often been stated in terms of maximizing utility, some have recently suggested that a less stringent and more attainable standard of expectations should be assumed. Fourth, utilitarians must indicate *what types of things are to be judged* in terms of their consequences. The three most common candidates here are acts, rules, and social policies. Finally, utilitarians must answer the question of *whom these are consequences for.* Clearly, they are not just the consequences for the individual agent—that would be ethical egoism. Do utilitarians take into account the consequences for all human beings or just for some subset, such as those in our own country? Do they take into account the consequences for future generations as well as the present one? Do they take into account the consequences for all sentient beings, animals as well as human beings; the natural environment as well as our constructed world; or just the human population?

Let's see how utilitarians have taken this basic insight about the moral significance of consequences and elaborated it into a formal theory of ethics.

DEFINING UTILITY

Utilitarians claim that the only thing that counts morally is whatever produces the greatest amount of *utility,* or the greatest overall positive consequences. Yet what is the proper yardstick of utility? Historically, utilitarians have taken pleasure and happiness as the measure of consequences. More recent versions of utilitarianism have turned either to higher ("ideal") goods or to preferences as the measure of consequences. Each of these four yardsticks has its strengths and its limitations.

Bentham and Pleasure

Originally, utilitarianism became influential with the work of Jeremy Bentham (1748–1832), who defined utility in terms of pleasure and pain. According to Bentham, we should act in such a way as to **maximize pleasure and minimize pain.** This position is known as **hedonistic utilitarianism.** Notice that it is very different from straightforward hedonism, which recommends maximizing one's own pleasure and minimizing one's own pain. Hedonistic utilitarianism recommends maximizing the overall amount of pleasure and minimizing the overall amount of pain.

Mill and Happiness

Bentham's philosophy quickly came under attack as the "pig's philosophy" because of what seemed to be its crude emphasis on sensual, bodily pleasures. John Stuart Mill (1806–1873), Bentham's godson, proposed a major

reformulation of the utilitarian position by arguing that **utility should be defined in terms of happiness rather than pleasure.** Mill's standard seemed to be a definite advance over Bentham's, for it was based on a higher standard than mere pleasure. This is called **eudaimonistic utilitarianism.** (The word *eudaimonistic* comes from the Greek word for happiness, *eudaimonia.*) In order to see why this is the case, let's consider some of the differences between pleasure and happiness as the standard of utility.

Pleasure versus Happiness

The differences between pleasure and happiness are significant. We tend to think of pleasure as being primarily bodily or sensual in character. Eating, drinking, and having sex come immediately to mind as paradigm cases of pleasure. Happiness, on the other hand, is usually less immediately tied down to the body. We might initially characterize it as belonging more to the mind or spirit than the body.

Second, pleasure generally seems to be of shorter duration than happiness. This stems from the nature of pleasure itself. Pleasure, at least in the eyes of many psychologists and philosophers, is the enjoyable feeling we experience when a state of deprivation is replaced by a state of satiation or fulfillment. In other words, pleasure is what we feel when we drink a nice cool glass of water when we are thirsty. Yet this gives us an insight into the reason why pleasures are short lived. Once we are satiated, we no longer experience the object as pleasurable. Once we are no longer thirsty, drinking water becomes less pleasurable. Happiness, on the other hand, seems to lie in the realization of certain goals, hopes, or plans for one's life. Insofar as these goals are intrinsically rewarding ones, we do not tire of them in the same way that we may tire of certain pleasures.

Third, happiness may encompass both pleasure and pain. Indeed, we could easily imagine persons saying that their life is happy but still acknowledging painful moments. A good example of this is a woman giving birth to a long-hoped-for child. She may experience quite a bit of pain during and after the birth, but she may still feel happy. Conversely, we can imagine someone experiencing pleasure but not feeling happy. Think of someone smoking crack, which directly stimulates the brain's pleasure center. That person might take pleasure from it as he inhales deeply, but he could be feeling very unhappy with his life, career, marriage, and so on.

Finally, there is more of an evaluative element in our notion of happiness than there is in our idea of pleasure. In reading the preceding example, many nonsmokers might have been repulsed at the idea of taking pleasure in smoking cigarettes, especially in the morning. Yet this is not a reason for doubting that some smokers do find pleasure in it. We may want to distinguish between good and bad pleasures, between harmless and harmful ones, but we do not doubt that the bad pleasures are still pleasures. With happiness, on the other hand, we build in an evaluative component. We are likely to question whether someone is genuinely happy in a way that we do not question whether they are genuinely feeling pleasure.

The problem with weighing consequences is that it is much easier to weigh pleasure than happiness or ideal goods, yet pleasure is the least suitable standard. *The closer we move toward a suitable standard of utility, the less able we are to subject it to quantification.*

Other Accounts of Utility

Pleasure and utility are not the only possible standards of utility, and the twentieth century saw attempts to redefine the standard of utility in terms of ideal goods such as freedom and knowledge and justice (G. E. Moore) and individual preferences (Kenneth Arrow). These versions, *ideal utilitarianism* and *preference utilitarianism* respectively, provide variations on the utilitarian theme. We can summarize these various versions of utilitarianism in the following way.

THE MEASURES OF UTILITY

Type of Utilitarianism	Standard of Utility	Number of Intrinsic Goods	Main Proponent
Hedonistic	Pleasure	One	Bentham
Eudaimonistic	Happiness	One	Mill
Ideal	Justice, Freedom, etc.	Many	Moore
Preference	Preference	None	Arrow

No single candidate has emerged as the sole choice among philosophers for the standard of utility. The disagreement among philosophers over this issue seems to reflect a wider disagreement in our own society. If consequences count, we have still to decide upon what yardstick to use in measuring them. The attraction of preference theory in this context is that it permits this multiplicity of standards, all of which are expressed as preferences.

Indeed, there may be a distinct advantage to allowing a *multiplicity* of different types of factors underlying utility. This is a type of pluralism within a specific moral theory. Utilitarianism has sometimes been criticized for being too narrow, for reducing all our considerations in life to a single axis of utility, usually either pleasure or happiness. There is much to be said for a fuller, suppler theory that permits us to recognize that consequences need to be measured according to several yardsticks. The difficulty with such a move, however, is that it then makes utilitarianism a more complex doctrine, one that is more difficult to apply in practice. Furthermore, utilitarians who go in this direction then need to specify the relationship among

the different kinds of yardsticks. When one alternative ranks high on the yardstick of happiness, for example, and another course of action is high on the scale of justice, which takes precedence? Finally, it threatens to rob utilitarianism of its chief advantage—namely, that it offers a clear method for calculating the morality of actions, rules, and social policies. Utilitarians who opt for a multiplicity of yardsticks must address themselves to questions such as these.

APPLYING THE MEASURE

Once we accept a standard of utility, we are still faced with the task of specifying how that standard is to apply to the world in which we live. This is a thorny but crucial issue for utilitarians because one of the major attractions of utilitarianism is that it promises precision in the moral life. If it can't deliver on this promise because it can't be applied precisely, then it loses its advantage over competing moral theories.

The Scale Metaphor

One of the things that makes utilitarianism attractive is the root metaphor that underlies much of its language about measuring consequences. Utilitarianism is grounded in a image with tremendous intuitive appeal to many of us: that of the scale. The very notion of weighing consequences presupposes that consequences are the kinds of things that can be placed on a scale. This metaphor pervades our everyday discourse about deciding among competing courses of action. Consider the types of things we often say.

- On *balance,* I'd rather go to the movies.
- When I *weigh* the alternatives, going to Hawaii looks best, even if it is hot at this time of year.
- Buying me an extra nice present *balanced* out the fact that they forgot to send it until a week after my birthday.
- Nothing can *outweigh* all the grief that the hit-and-run accident caused us.
- Only the death penalty can right the *scales* of justice.
- This is a *weighty* choice with *heavy* consequences.

Some of these scale metaphors have specifically monetary overtones, as though things were weighed in terms of their dollar value.

- I'm going to *pay him back* for all the grief he's given me over the years.
- How can I ever *repay* you for your kindness?
- That was a *costly* mistake.
- I'll be forever in your *debt.*

In these expressions, we see the way in which money metaphorically plays the role of the measure in terms of which consequences are assessed.

In order to weigh consequences, the utilitarian needs some kind of measure in terms of which utility is determined. We have examined various candidates for this yardstick: pleasure, happiness, ideals, and preferences. But we also need some way of marking off the units to be measured. Scales are often marked off in terms of ounces or grams. Yardsticks are usually marked off in terms of inches. How do we mark off units of utility? One of the ways of doing this is by assigning cardinal numbers (1, 2, 3, etc.) to pleasure or happiness, for example. In the following section, we will consider this approach by using the stipulative concepts of hedons and dolors. The other way is to assign ordinal (1st, 2nd, 3rd, etc.) or coordinal (comparative) rankings. This is the approach that preference utilitarianism takes. Let's look at both.

Hedons and Dolors

Once they agree upon the kind of yardstick, utilitarians must arrive at some consensus about the way in which the individual units on the measuring stick are to be marked off. Utilitarians sometimes refer to units of pleasure or happiness as **hedons** and units of displeasure or suffering or unhappiness as **dolors.** (The word *hedon* comes from the Greek word for *pleasure;* this is the same root that *hedonism* comes from. The word *dolor* comes from the Latin word *dolor,* which means "pain.") Particular things, at least in textbook examples, are then often assigned some number. A good corned beef sandwich may, for example, be 5 hedons, while a pleasant vacation to the Bahamas may be 6,000 hedons. A visit to the dentist could be 100 dolors, and the death of a close friend several thousand dolors.

Such a system may at first sound artificial, but utilitarians would argue that this apparent artificiality is not a serious problem. While there may be no absolute scale in which going to the dentist is a 100, it may be the case that having a close friend die is roughly twenty or thirty times worse than going to the dentist. It is this relationship of relative suffering (or pleasure) that the utilitarian seeks to capture in assigning numerical values to various consequences. This is, the utilitarian further argues, something that we do quite naturally in our everyday lives. The **utilitarian calculus** is but a refinement and formalization of that everyday activity of assigning relative values to various occurrences, of ranking them in relation to each other according to the amount of pleasure or pain they yield. But as soon as cardinal utilitarians admit this, they are on their way to becoming preference utilitarians.

The Decision Procedure

How do utilitarians go about deciding upon the moral worth of an action, granting that the consequences of an action can be specified in terms of hedons and dolors? They claim that in any given situation we must to the best of our ability (1) determine the consequences of the various courses of action

open to us, (2) specify the hedons and dolors associated with each alternative, and then (3) perform that course of action that results in the greatest total amount of pleasure (i.e., of hedons minus dolors). Imagine, for example, that you are a utilitarian in the process of deciding between two pieces of proposed legislation about medical aid for the elderly. There are three possibilities open to you: to vote for a bill to reduce medical aid, to vote in favor of a bill that would increase such aid, or to vote against both and thus effectively vote in favor of keeping things the same. Reducing medical benefits for the elderly may result in 10 hedons apiece for 100 million people and 200 dolors for 20 million people, resulting in an overall utility of 3 billion dolors. Keeping benefits the same may result in 20 hedons apiece for 20 million people and 3 dolors for 100 million people, with a total overall utility of 2 billion, 600 million hedons. Finally, increasing benefits may result in 90 hedons apiece for 20 million people and 20 dolors for 100 million people, with a total overall utility of 200 million hedons. Thus, from a utilitarian point of view, we would be obligated to keep the benefits the same, since the other two possible courses of action both have a lesser overall utility.

How Much Utility Is Enough?

One of the difficulties that utilitarians have faced centers around the question of how much utility we are obligated to produce. The usual answer has been that we ought to do whatever produces the greatest overall amount of utility. When we compare competing courses of action, we should choose the best one, the one that maximizes utility.

Maximizing Utility

It is easy to understand the initial plausibility of this answer. Imagine that we are considering whether to pass a particular piece of social legislation. We naturally consider how we can produce the greatest amount of good. If we are weighing the alternatives impartially, there is no attraction toward anything else, no pull toward doing less than the best. If some particular special interest group desired that we do less than the best, their wishes would be factored into the utilitarian equation along with everyone else's. There would, however, be no reason for giving them special weight. If we thought they deserved more weight, then presumably it would be for some reason that would be recognizable within the utilitarian framework. It is only when we have partial interests of our own that there is a desire to do less than the best.

Yet the picture changes rapidly when we consider what it would be like to be making decisions about our own personal lives in the same way. Here personal desires and wants have a much more prominent role. Yet traditional versions of utilitarianism seem to demand that we do the maximum, even in these more personal situations. The utilitarian always tries to produce the greatest overall amount of utility. This is in sharp contrast, for example, to Kantian ethics, which we will discuss in detail in the next chapter. Kant's position just states negatively that a particular action is morally forbidden, while the utilitarian tells us positively that we must choose the specific course of action which maximizes utility. Thus utilitarianism is an extremely demanding moral doctrine, since it demands that we sacrifice our own pleasure, happiness, or preference satisfaction for the greater good—that is, for social utility.

Supererogation

Philosophers have noted an interesting implication of the utilitarian view that we must always try to maximize utility. In many moral philosophies, it is often possible to act in a **supererogatory** fashion—that is, to go beyond the demands of duty and thus do something exceptionally meritorious. For the utilitarian, however, this is impossible. One is always obligated to do the thing that yields the greatest amount of utility, and it is precisely this obligation that constitutes duty. Thus, there is nothing above the call of duty. For the utilitarian, there is no room for supererogatory actions, for duty is so demanding that nothing above it is greater.

This is a cause for concern among some philosophers, for they believe that a moral theory that has no room for supererogation must be mistaken in some important way. Our everyday moral intuitions tell us that sometimes people do something that is above the call of duty. But if utilitarianism is correct, then this is impossible. Duty always calls for the maximum, so it is impossible to do anything above its call.

Consequences of What?

The utilitarian maintains that we ought to prefer whatever produces the greatest overall utility, and this is determined by weighing the consequences. But the consequences *of what*? Utilitarians have given at least three different answers to this question, answers that are not necessarily mutually exclusive: acts, rules, and practices.

Act Utilitarianism

The first and most common version of utilitarianism says that **we should look at the consequences of each individual action in attempting to determine its moral worth.** This position, which is called **act utilitarianism,** maintains that we should always perform that *action* that will maximize utility,